Package 'rvest'

March 9, 2021

Title Easily Harvest (Scrape) Web Pages

Version 1.0.0

Description Wrappers around the 'xml2' and 'httr' packages to
make it easy to download, then manipulate, HTML and XML.

License MIT + file LICENSE

URL <https://rvest.tidyverse.org/>, <https://github.com/tidyverse/rvest>

BugReports <https://github.com/tidyverse/rvest/issues>

Depends R (>= 3.2)

Imports httr (>= 0.5), lifecycle (>= 1.0.0), magrittr, rlang (>=
0.4.10), selectr, tibble, xml2 (>= 1.3)

Suggests covr, glue, knitr, readr, rmarkdown, repurrrsive, spelling,
stringi (>= 0.3.1), testthat (>= 3.0.2), webfakes

VignetteBuilder knitr

Config/testthat/edition 3

Encoding UTF-8

Language en-US

LazyData true

RoxygenNote 7.1.1

NeedsCompilation no

Author Hadley Wickham [aut, cre],
RStudio [cph]

Maintainer Hadley Wickham <hadley@rstudio.com>

Repository CRAN

Date/Publication 2021-03-09 16:50:02 UTC

R topics documented:

---

html_attr   *Get element attributes*

Description

html_attr() gets a single attribute; html_attr() gets all attributes.

Usage

```
html_attr(x, name, default = NA_character_)

html_attrs(x)
```

Arguments

x   A document (from read_html()), node set (from html_elements()), node
            (from html_element()), or session (from session()).

name   Name of attribute to retrieve.

default   A string used as a default value when the attribute does not exist in every ele-
            ment.

Value

A character vector (for html_attr()) or list (html_attrs()) the same length as x.

Examples

```
url <- "https://en.wikipedia.org/w/index.php?title=The_Lego_Movie&oldid=998422565"
html <- read_html(url)

cast <- html_elements(html, "tr:nth-child(8) .plainlist a")
cast %>% html_text2()
cast %>% html_attrs()
cast %>% html_attr("href")
```

```
# If needed, use url_absolute() to convert to complete urls
url_absolute(html_attr(cast, "href"), url)
```

---

html_children   *Get element children*

Description

Get element children

Usage

```
html_children(x)
```

Arguments

x   A document (from read_html()), node set (from html_elements()), node
                (from html_element()), or session (from session()).

Examples

```
html <- minimal_html("<ul><li>1<li>2<li>3</ul>")
ul <- html_elements(html, "ul")
html_children(ul)

html <- minimal_html("<p>Hello <b>Hadley</b><i>!</i>")
p <- html_elements(html, "p")
html_children(p)
```

---

html_element   *Select elements from an HTML document*

Description

html_element() and html_elements() find HTML element using CSS selectors or XPath expressions. CSS selectors are particularly useful in conjunction with https://selectorgadget.com/, which makes it very easy to discover the selector you need.

Usage

```
html_element(x, css, xpath)

html_elements(x, css, xpath)
```

Arguments

x	Either a document, a node set or a single node.
css, xpath	Elements to select. Supply one of `css` or `xpath` depending on whether you want to use a CSS selector or XPath 1.0 expression.

Value

`html_element()` returns a nodeset the same length as the input. `html_elements()` flattens the output so there's no direct way to map the output to the input.

CSS selector support

CSS selectors are translated to XPath selectors by the **selectr** package, which is a port of the python **cssselect** library, https://pythonhosted.org/cssselect/.

It implements the majority of CSS3 selectors, as described in http://www.w3.org/TR/2011/REC-css3-selectors-20110929/. The exceptions are listed below:

- Pseudo selectors that require interactivity are ignored: :hover, :active, :focus, :target, :visited.
- The following pseudo classes don't work with the wild card element, *: *:first-of-type, *:last-of-type, *:nth-of-type, *:nth-last-of-type, *:only-of-type
- It supports :contains(text)
- You can use !=, [foo!=bar] is the same as :not([foo=bar])
- :not() accepts a sequence of simple selectors, not just a single simple selector.

Examples

```
html <- minimal_html("
  <h1>This is a heading</h1>
  <p id='first'>This is a paragraph</p>
  <p class='important'>This is an important paragraph</p>
")

html %>% html_element("h1")
html %>% html_elements("p")
html %>% html_elements(".important")
html %>% html_elements("#first")

# html_element() vs html_elements() -------------------------------------
html <- minimal_html("
  <ul>
    <li><b>C-3PO</b> is a <i>droid</i> that weighs <span class='weight'>167 kg</span></li>
    <li><b>R2-D2</b> is a <i>droid</i> that weighs <span class='weight'>96 kg</span></li>
    <li><b>Yoda</b> weighs <span class='weight'>66 kg</span></li>
    <li><b>R4-P17</b> is a <i>droid</i></li>
  </ul>
")
li <- html %>% html_elements("li")

# When applied to a node set, html_elements() returns all matching elements
```

```
# beneath any of the inputs, flattening results into a new node set.
li %>% html_elements("i")

# When applied to a node set, html_element() always returns a vector the
# same length as the input, using a "missing" element where needed.
li %>% html_element("i")
# and html_text() and html_attr() will return NA
li %>% html_element("i") %>% html_text2()
li %>% html_element("span") %>% html_attr("class")
```

---

html_encoding_guess   *Guess faulty character encoding*

Description

html_encoding_guess() helps you handle web pages that declare an incorrect encoding. Use html_encoding_guess() to generate a list of possible encodings, then try each out by using encoding argument of read_html(). html_encoding_guess() replaces the deprecated guess_encoding().

Usage

```
html_encoding_guess(x)
```

Arguments

x   A character vector.

Examples

```
# A file with bad encoding included in the package
path <- system.file("html-ex", "bad-encoding.html", package = "rvest")
x <- read_html(path)
x %>% html_elements("p") %>% html_text()

html_encoding_guess(x)
# Two valid encodings, only one of which is correct
read_html(path, encoding = "ISO-8859-1") %>% html_elements("p") %>% html_text()
read_html(path, encoding = "ISO-8859-2") %>% html_elements("p") %>% html_text()
```

---

html_form   *Parse forms and set values*

Description

Use html_form() to extract a form, set values with html_form_set(), and submit it with html_form_submit().

Usage

```
html_form(x, base_url = NULL)

html_form_set(form, ...)

html_form_submit(form, submit = NULL)
```

Arguments

x	A document (from [read_html()](#), node set (from [html_elements()](#)), node (from [html_element()](#)), or session (from [session()](#)).
base_url	Base url of underlying HTML document. The default, NULL, uses the url of the HTML document underlying x.
form	A form
...	<[dynamic-dots](#)> Name-value pairs giving fields to modify. Provide a character vector to set multiple checkboxes in a set or select multiple values from a multi-select.
submit	Which button should be used to submit the form? • NULL, the default, uses the first button. • A string selects a button by its name. • A number selects a button using its relative position.

Value

- html_form() returns as S3 object with class rvest_form when applied to a single element. It returns a list of rvest_form objects when applied to multiple elements or a document.
- html_form_set() returns an rvest_form object.
- html_form_submit() submits the form, returning an httr response which can be parsed with [read_html()](#).

See Also

HTML 4.01 form specification: <http://www.w3.org/TR/html401/interact/forms.html>

Examples

```
html <- read_html("http://www.google.com")
search <- html_form(html)[[1]]

search <- search %>% html_form_set(q = "My little pony", hl = "fr")

# Or if you have a list of values, use !!!
vals <- list(q = "web scraping", hl = "en")
search <- search %>% html_form_set(!!!vals)

# To submit and get result:
## Not run:
resp <- html_form_submit(search)
read_html(resp)

## End(Not run)
```

---

html_name	*Get element name*

Description

Get element name

Usage

```
html_name(x)
```

Arguments

x A document (from [read_html()](), node set (from [html_elements()]()), node (from [html_element()]()), or session (from [session()]()).

Value

A character vector the same length as x

Examples

```
url <- "https://rvest.tidyverse.org/articles/starwars.html"
html <- read_html(url)

html %>%
  html_element("div") %>%
  html_children() %>%
  html_name()
```

---

html_table	*Parse an html table into a data frame*

Description

The algorithm mimics what a browser does, but repeats the values of merged cells in every cell that cover.

Usage

```
html_table(
  x,
  header = NA,
  trim = TRUE,
  fill = deprecated(),
  dec = ".",
  na.strings = "NA",
  convert = TRUE
)
```

Arguments

x	A document (from read_html()), node set (from html_elements()), node (from html_element()), or session (from session()).
header	Use first row as header? If NA, will use first row if it consists of <th> tags.
	If TRUE, column names are left exactly as they are in the source document, which may require post-processing to generate a valid data frame.
trim	Remove leading and trailing whitespace within each cell?
fill	Deprecated - missing cells in tables are now always automatically filled with NA.
dec	The character used as decimal place marker.
na.strings	Character vector of values that will be converted to NA if convert is TRUE.
convert	If TRUE, will run type.convert() to interpret texts as integer, double, or NA.

Value

When applied to a single element, html_table() returns a single tibble. When applied to multiple elements or a document, html_table() returns a list of tibbles.

Examples

```
sample1 <- minimal_html("<table>
  <tr><th>Col A</th><th>Col B</th></tr>
  <tr><td>1</td><td>x</td></tr>
  <tr><td>4</td><td>y</td></tr>
  <tr><td>10</td><td>z</td></tr>
</table>")
```

```
sample1 %>%
  html_element("table") %>%
  html_table()

# Values in merged cells will be duplicated
sample2 <- minimal_html("<table>
  <tr><th>A</th><th>B</th><th>C</th></tr>
  <tr><td>1</td><td>2</td><td>3</td></tr>
  <tr><td colspan='2'>4</td><td>5</td></tr>
  <tr><td>6</td><td colspan='2'>7</td></tr>
</table>")
sample2 %>%
  html_element("table") %>%
  html_table()

# If a row is missing cells, they'll be filled with NAs
sample3 <- minimal_html("<table>
  <tr><th>A</th><th>B</th><th>C</th></tr>
  <tr><td colspan='2'>1</td><td>2</td></tr>
  <tr><td colspan='2'>3</td></tr>
  <tr><td>4</td></tr>
</table>")
sample3 %>%
  html_element("table") %>%
  html_table()
```

---

html_text   *Get element text*

Description

There are two ways to retrieve text from a element: `html_text()` and `html_text2()`. `html_text()` is a thin wrapper around `xml2::xml_text()` which returns just the raw underlying text. `html_text2()` simulates how text looks in a browser, using an approach inspired by JavaScript's innerText(). Roughly speaking, it converts `<br />` to `"\n"`, adds blank lines around `<p>` tags, and lightly formats tabular data.

`html_text2()` is usually what you want, but it is much slower than `html_text()` so for simple applications where performance is important you may want to use `html_text()` instead.

Usage

```
html_text(x, trim = FALSE)

html_text2(x, preserve_nbsp = FALSE)
```

Arguments

x	A document, node, or node set.
trim	If TRUE will trim leading and trailing spaces.

preserve_nbsp   Should non-breaking spaces be preserved? By default, `html_text2()` converts to ordinary spaces to ease further computation. When `preserve_nbsp` is TRUE,   will appear in strings as `"\ua0"`. This often causes confusion because it prints the same way as `" "`.

Value

A character vector the same length as x

Examples

```
# To understand the difference between html_text() and html_text2()
# take the following html:

html <- minimal_html(
  "<p>This is a paragraph.
    This another sentence.<br>This should start on a new line"
)

# html_text() returns the raw underlying text, which includes whitespace
# that would be ignored by a browser, and ignores the <br>
html %>% html_element("p") %>% html_text() %>% writeLines()

# html_text2() simulates what a browser would display. Non-significant
# whitespace is collapsed, and <br> is turned into a line break
html %>% html_element("p") %>% html_text2() %>% writeLines()

# By default, html_text2() also converts non-breaking spaces to regular
# spaces:
html <- minimal_html("<p>x y</p>")
x1 <- html %>% html_element("p") %>% html_text()
x2 <- html %>% html_element("p") %>% html_text2()

# When printed, non-breaking spaces look exactly like regular spaces
x1
x2
# But aren't actually the same:
x1 == x2
# Which you can confirm by looking at their underlying binary
# representaion:
charToRaw(x1)
charToRaw(x2)
```

---

session   *Simulate a session in web browser*

Description

This set of functions allows you to simulate a user interacting with a website, using forms and navigating from page to page.

- Create a session with `session(url)`
- Navigate to a specified url with `session_jump_to()`, or follow a link on the page with `session_follow_link()`.
- Submit an [html_form](#) with `session_submit()`.
- View the history with `session_history()` and navigate back and forward with `session_back()` and `session_forward()`.
- Extract page contents with [html_element()](#) and [html_elements()](#), or get the complete HTML document with [read_html()](#).
- Inspect the HTTP response with [httr::cookies()](#), [httr::headers()](#), and [httr::status_code()](#).

Usage

```
session(url, ...)

is.session(x)

session_jump_to(x, url, ...)

session_follow_link(x, i, css, xpath, ...)

session_back(x)

session_forward(x)

session_history(x)

session_submit(x, form, submit = NULL, ...)
```

Arguments

`url`	A URL, either relative or absolute, to navigate to.
`...`	Any additional httr config to use throughout the session.
`x`	A session.
`i`	A integer to select the ith link or a string to match the first link containing that text (case sensitive).
`css`	Elements to select. Supply one of `css` or `xpath` depending on whether you want to use a CSS selector or XPath 1.0 expression.
`xpath`	Elements to select. Supply one of `css` or `xpath` depending on whether you want to use a CSS selector or XPath 1.0 expression.
`form`	An [html_form](#) to submit
`submit`	Which button should be used to submit the form?

- NULL, the default, uses the first button.
- A string selects a button by its name.
- A number selects a button using its relative position.

Examples

```
s <- session("http://hadley.nz")
s %>%
  session_jump_to("hadley-wickham.jpg") %>%
  session_jump_to("/") %>%
  session_history()

s %>%
  session_jump_to("hadley-wickham.jpg") %>%
  session_back() %>%
  session_history()


s %>%
  session_follow_link(css = "p a") %>%
  html_elements("p")
```

Index